

# CS 2710 Foundations of AI

## Lecture 8

### Adversarial search

**Milos Hauskrecht**

[milos@cs.pitt.edu](mailto:milos@cs.pitt.edu)

5329 Sennott Square

---

CS 2710 Foundations of AI

### Game search

- Game-playing programs developed by AI researchers since the beginning of the modern AI era
  - Programs playing chess, checkers, etc (1950s)
- **Specifics of the game search:**
  - Sequences of player's decisions **we control**
  - Decisions of other player(s) **we do not control**
- **Contingency problem:** many possible opponent's moves must be “covered” by the solution

Opponent's behavior introduces an uncertainty in to the game

  - We do not know exactly what the response is going to be
- **Rational opponent** – maximizes it own **utility (payoff) function**

---

CS 2710 Foundations of AI

## Types of game problems

- Types of game problems:
  - **Adversarial games:**
    - win of one player is a loss of the other
  - **Cooperative games:**
    - players have common interests and utility function
  - A spectrum of game problems in between the two:

Adversarial games

Fully cooperative games




**we focus on adversarial games only!!**

CS 2710 Foundations of AI

## Example of an adversarial 2 person game: Tic-tac-toe

- Player 1 (x) moves first



CS 2710 Foundations of AI

## Game search problem

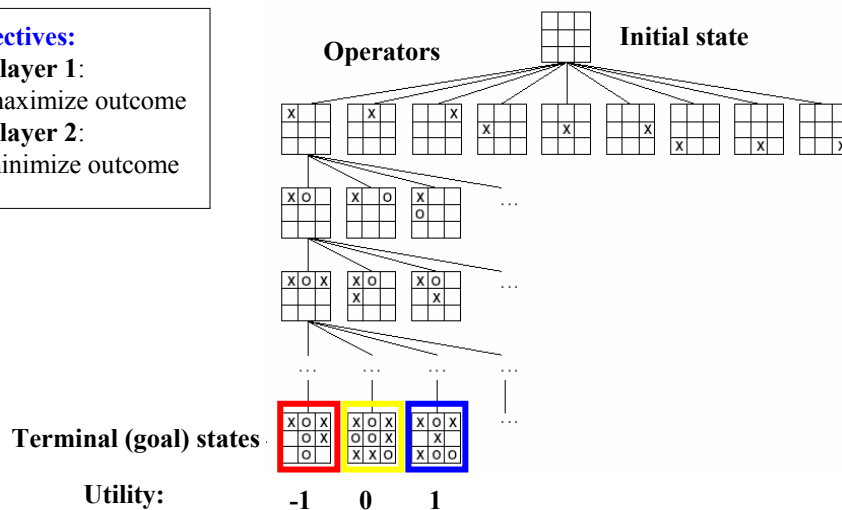
- **Game problem formulation:**
  - **Initial state:** initial board position + info whose move it is
  - **Operators:** legal moves a player can make
  - **Goal (terminal test):** determines when the game is over
  - **Utility (payoff) function:** measures the outcome of the game and its desirability
- **Search objective:**
  - find the sequence of player's decisions (moves) maximizing its utility (payoff)
  - Consider the opponent's moves and their utility

CS 2710 Foundations of AI

## Game problem formulation (Tic-tac-toe)

### Objectives:

- **Player 1:** maximize outcome
- **Player 2:** minimize outcome

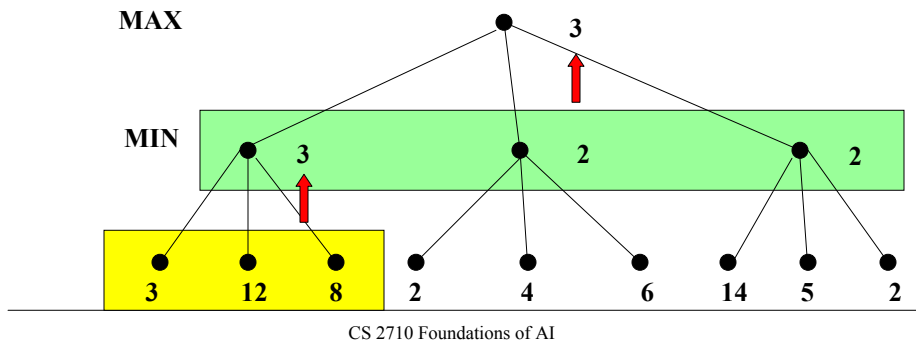


CS 2710 Foundations of AI


## Minimax algorithm

How to deal with the contingency problem?


- Assuming that the opponent is rational and always optimizes its behavior (opposite to us) we consider **the best opponent's response**
- Then the **minimax algorithm** determines the best move



## Minimax algorithm. Example




## Minimax algorithm. Example




CS 2710 Foundations of AI

## Minimax algorithm. Example




CS 2710 Foundations of AI

## Minimax algorithm. Example



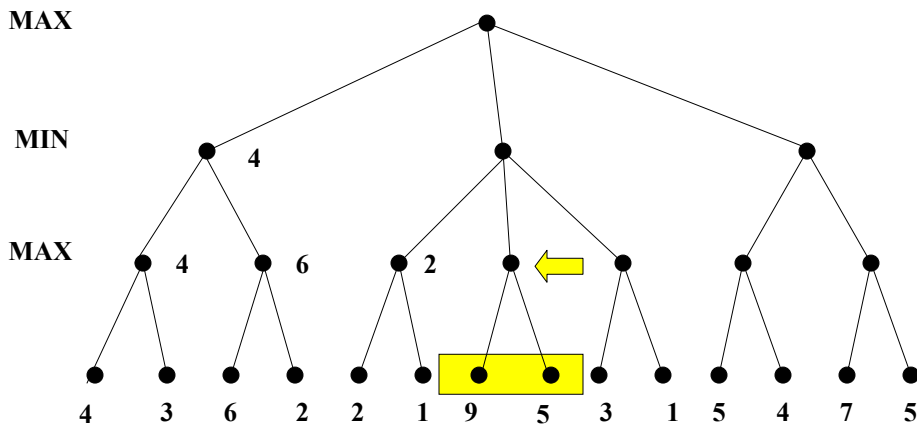
CS 2710 Foundations of AI

## Minimax algorithm. Example




CS 2710 Foundations of AI

## Minimax algorithm. Example




CS 2710 Foundations of AI

## Minimax algorithm. Example




CS 2710 Foundations of AI

## Minimax algorithm. Example



CS 2710 Foundations of AI

## Minimax algorithm. Example



CS 2710 Foundations of AI



# Minimax algorithm

```

function MINIMAX-DECISION(game) returns an operator
  for each op in OPERATORS[game] do
    VALUE[op]  $\leftarrow$  MINIMAX-VALUE(APPLY(op, game), game)
  end
  return the op with the highest VALUE[op]
  
```


```

function MINIMAX-VALUE(state, game) returns a utility value
  if TERMINAL-TEST[game](state) then
    return UTILITY[game](state)
  else if MAX is to move in state then
    return the highest MINIMAX-VALUE of SUCCESSORS(state)
  else
    return the lowest MINIMAX-VALUE of SUCCESSORS(state)
  
```

CS 2710 Foundations of AI

## Complexity of the minimax algorithm


- We need to explore the complete game tree before making the decision



CS 2710 Foundations of AI

## Complexity of the minimax algorithm

- We need to explore the complete game tree before making the decision



- Impossible for large games
  - Chess: 35 operators, game can have 50 or more moves

## Solution to the complexity problem

### Two solutions:

1. **Dynamic pruning of redundant branches** of the search tree
  - identify a provably suboptimal branch of the search tree before it is fully explored
  - Eliminate the suboptimal branch


### Procedure: **Alpha-Beta pruning**

2. **Early cutoff of the search tree**
  - uses imperfect minimax value estimate of non-terminal states (positions)


## Alpha beta pruning

- Some branches will never be played by rational players since they include sub-optimal decisions (for either player)

## Alpha beta pruning. Example




## Alpha beta pruning. Example




CS 2710 Foundations of AI

## Alpha beta pruning. Example



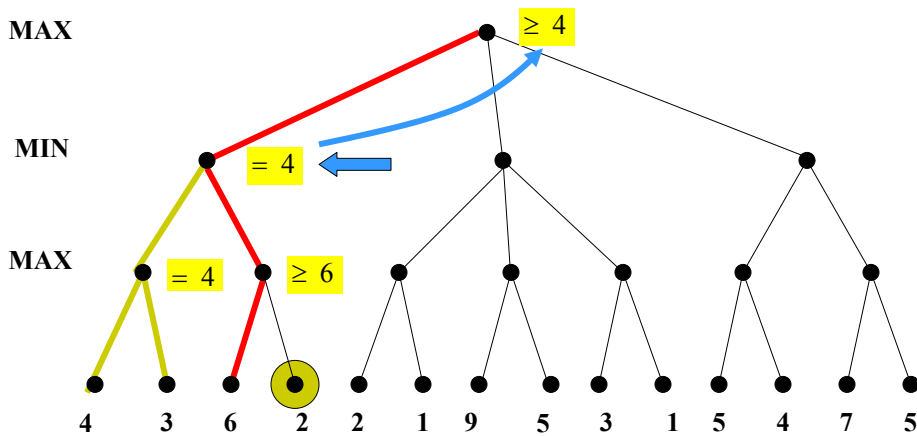
CS 2710 Foundations of AI

## Alpha beta pruning. Example




CS 2710 Foundations of AI

## Alpha beta pruning. Example



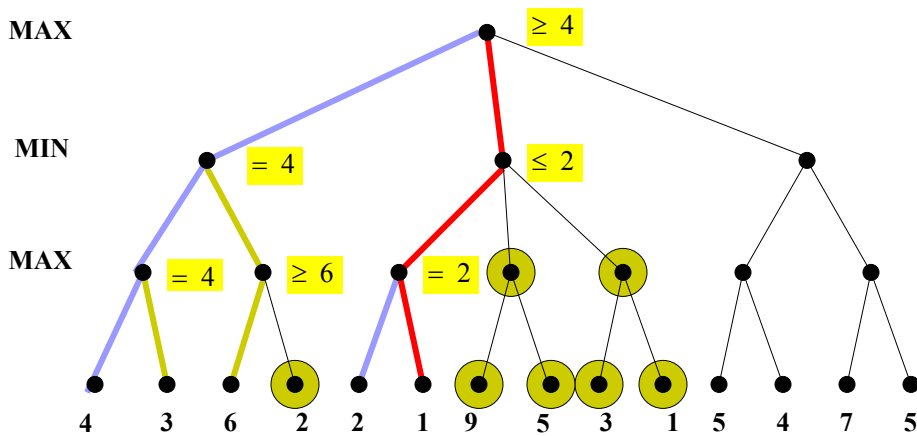
CS 2710 Foundations of AI

## Alpha beta pruning. Example




CS 2710 Foundations of AI

## Alpha beta pruning. Example



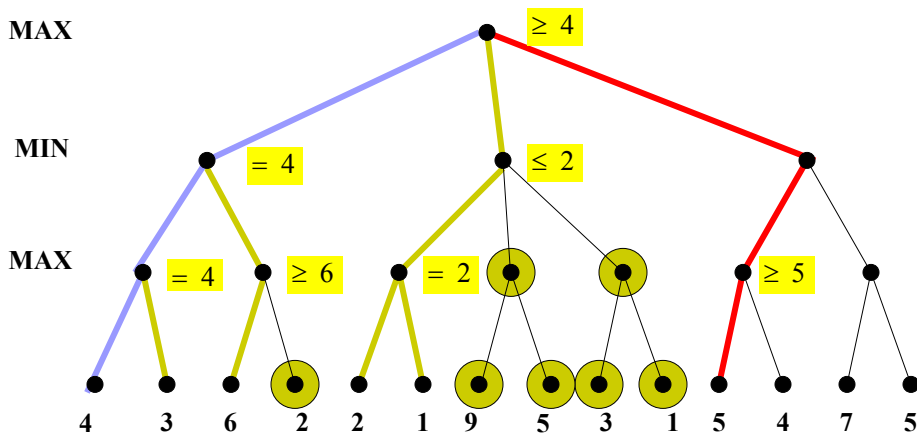
CS 2710 Foundations of AI

## Alpha beta pruning. Example




CS 2710 Foundations of AI

## Alpha beta pruning. Example



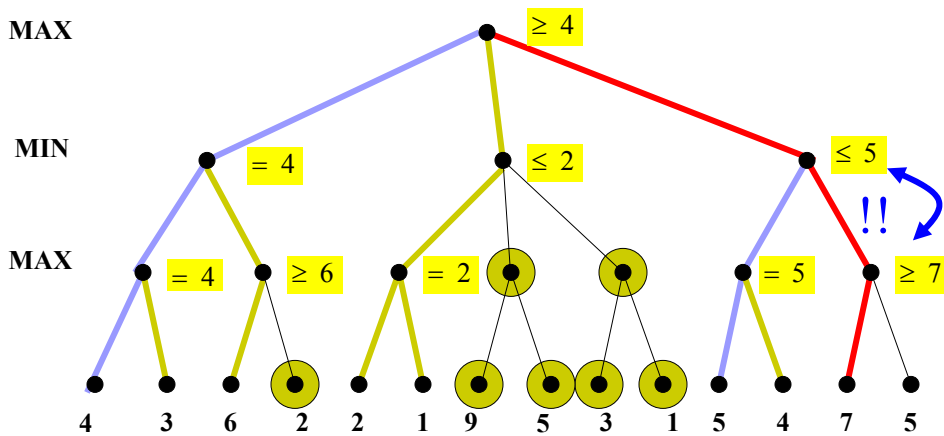
CS 2710 Foundations of AI

## Alpha beta pruning. Example



CS 2710 Foundations of AI


## Alpha beta pruning. Example



CS 2710 Foundations of AI




## Alpha beta pruning. Example



CS 2710 Foundations of AI

## Alpha beta pruning. Example



CS 2710 Foundations of AI

## Alpha-Beta pruning

**function** MAX-VALUE(*state*, *game*,  $\alpha$ ,  $\beta$ ) **returns** the minimax value of *state*

**inputs:** *state*, current state in game

*game*, game description

$\alpha$ , the best score for MAX along the path to *state*

$\beta$ , the best score for MIN along the path to *state*

**if** GOAL-TEST(*state*) **then return** EVAL(*state*)

**for each** *s* **in** SUCCESSORS(*state*) **do**

$\alpha \leftarrow \text{MAX}(\alpha, \text{MIN-VALUE}(s, \text{game}, \alpha, \beta))$

**if**  $\alpha \geq \beta$  **then return**  $\beta$

**end**

**return**  $\alpha$

---

**function** MIN-VALUE(*state*, *game*,  $\alpha$ ,  $\beta$ ) **returns** the minimax value of *state*

**if** GOAL-TEST(*state*) **then return** EVAL(*state*)

**for each** *s* **in** SUCCESSORS(*state*) **do**

$\beta \leftarrow \text{MIN}(\beta, \text{MAX-VALUE}(s, \text{game}, \alpha, \beta))$

**if**  $\beta \leq \alpha$  **then return**  $\alpha$

**end**


**return**  $\beta$

CS 2710 Foundations of AI

## Using minimax value estimates

- Idea:**

- Cutoff the search tree before the terminal state is reached
- Use imperfect estimate of the minimax value at the leaves
  - Evaluation function



CS 2710 Foundations of AI

## Design of evaluation functions

- **Heuristic estimate** of the value for a sub-tree
- **Examples of a heuristic functions:**
  - **Material advantage in chess, checkers**
    - Gives a value to every piece on the board, its position and combines them
  - More general **feature-based evaluation function**
    - Typically a linear evaluation function:

$$f(s) = f_1(s)w_1 + f_2(s)w_2 + \dots f_k(s)w_k$$

$f_i(s)$  - a feature of a state  $s$

$w_i$  - feature weight

## Further extensions to real games

- Restricted set of moves to be considered under **the cutoff level** to reduce branching and improve the evaluation function
  - E.g., consider only the capture moves in chess

